How the intelligence system is connected

A local-first pipeline connects configured public sources, scheduled collectors, a shared SQLite knowledge base, deterministic analysis, two AI layers, and the browser dashboard.

Design principle: evidence and workflow state stay local

The local Python application owns source collection, scoring, storage, scheduling, analysis, and delivery. External AI receives bounded context for synthesis or review; it does not own the database or the run schedule.

1 COLLECT — Source inputs

- **Public web + RSS**
  Market, retail, product, hiring, partnership, research, and news sources.
- **Focused source sets**
  44 market sources, 11 competitor sources, 26 partner sources, and 7 sales-lead sources.
- **Watch + inbox**
  Saved competitors, keywords, web pages, images, and the research inbox.

2 CONTROL — Local core

- **Scheduler threads**
  Market, competitor, partner, sales, research, sizing, and reporting runners.
- **Collector controls**
  Robots checks, polite fetching, source rotation, budgets, normalization, scoring, and deduplication.
- **Shared SQLite store**
  Signals, leads, competitors, sources, dossiers, drafts, snapshots, and run history.

3 ANALYZE — Intelligence

- **Deterministic engines**
  Trends, competitive gaps, opportunities, market sizing, kiosk marketing, profiles, and reports.
- **Gemini layer**
  Summaries, company research, section copilots, JSON designs, and source-grounded answers.
- **AgentCrew layer**
  Ten named crews review bounded workflow payloads and return structured judgment.

4 DELIVER — Experiences

- **Dashboard**
  Nine workspaces with highlights, expandable evidence, status controls, and section-specific...
- **Outputs**
  Leadership digest, scheduled email reports, research drafts, company dossiers, and generated PDFs.
- **Admin surfaces**
  Crew library, crew designer, status, run controls, and evidence/source management.

Two distinct AI connections

DIRECT GEMINI

Interactive synthesis and copilot answers

The server builds section context, sends it to Gemini, and returns cited answers. Browser clients never receive service credentials.

AGENTCREW API

Workflow review and multi-agent judgment

The local runtime submits bounded run payloads to named crews, polls for structured output, and records a safe local fallback when remote execution fails.

Primary implementation: cli.py, scheduler.py, storage.py, dashboard.py, llm.py, agentcrew.py, multi_agent_jobs.py

What runs, when, and where data moves

Each scheduled runner follows the same control loop: collect within a budget, write normalized records locally, compute analysis, optionally request crew review, then expose the result through APIs and the dashboard.

The local run loop

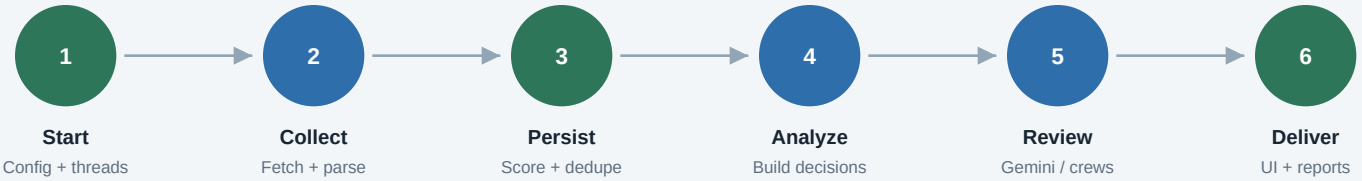

Configured background schedules

360m — Market intelligence
Three collection passes, SQLite write, digest, then market crew review.

72m — Competitor intelligence
Rotating feed sources, watchlist/browser/profile work, then crew review.

90m — Partners + resellers
Rotating partner sources, contact evidence, qualification, and crew review.

120m — Sales leads
Business triggers, product-fit candidates, evidence profiles, and crew review.

60m — Research automation
Saved searches, collections, drafts, dossiers, inbox, and crew review.

Wkly — Sizing + reporting
Market-sizing evidence reviews and scheduled leadership email preparation.

RUNTIME SNAPSHOT - 04 AUG 2026

290	14	10/10	No
stored signals	signals added today	crews registered	remote execution verified

Operational truth

FRESHNESS + AVAILABILITY

- Schedulers run only while the Python server remains open and the host machine is awake.
- The UI polls runner status every 30 seconds; executive metrics refresh on page load or when the user clicks Refresh.
- Robots rules, 404 pages, source limits, and missing browser tooling can reduce coverage without stopping the system.

SECURITY + RESILIENCY

- API credentials stay server-side and are loaded from environment settings; the browser calls local endpoints only.
- SQLite remains the source of record. AgentCrew is configured with database_writeback=false.
- If remote crew execution fails, collection, persistence, analysis, and delivery complete through the local safeguard path.

How the crew layer is connected through the API

AgentCrew is an external review service attached to local workflow checkpoints. It adds role-based judgment and structured review, while the local application keeps deterministic control of data, schedules, and delivery.

LOCAL INTELLIGENCE RUNTIME

Builds a bounded JSON payload, chooses the named system crew, submits the job, polls status, and records the result.

HTTPS + JSON →

AGENTCREW SERVICE

Authenticates the tenant, stores crew definitions, executes agents and tools, then returns a normalized job result.

API contract used by the integration

METHOD	REMOTE ROUTE	HOW THE LOCAL CLIENT USES IT
POST	`/api/v1/login`	Username/password login; cache returned Bearer token.
GET	`/api/v1/crew/tools`	Discover tools available to crew definitions.
GET	`/api/v1/crew?tenant=...`	Resolve existing crews and stable crew IDs by tenant.
PUT	`/api/v1/crew`	Create a missing system crew from the Gemini-designed definition.
POST	`/api/v1/crews`	Submit crew_id, query, tenant, mode, token and temperature limits.
PATCH	`/api/v1/crews?job_id=...`	Poll every 2 seconds until complete, failed, or 180-second timeout.

Execution sequence

1. **Configure** — Base URL, token or login, tenant, default Gemini model
2. **Register** — Find or create deterministic mi_<job>_crew definitions
3. **Trigger** — Collector or analysis checkpoint supplies a bounded payload
4. **Execute** — POST job, receive job_id, then poll status
5. **Normalize** — Extract JSON structured_output and preserve evidence boundaries
6. **Safeguard** — On any error, record local_fallback and keep local work active

Ten system crews mapped to workflow checkpoints

- Market intelligence
- General market trends
- Research automation
- Sales lead discovery
- Partner discovery
- Competitor intelligence
- Opportunity analysis
- In-store marketing
- Market sizing
- Leadership reporting

LOCAL SYSTEM OWNS

Collection, validation rules, SQLite writes, scheduling, analysis engines, UI, report delivery, and fallback completion.

AGENTCREW OWNS

Crew definitions, agent/tool execution, role-based review, and structured judgment returned to the local runtime.

CURRENT STATE 10/10 crews registered; remote execution is not yet verified. Recent remote jobs returned Not Found, so local_safeguard completed the workflow.